



Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

LEARNER GUIDE

For

Generative AI for SEO

Course Code: C11

Conducted by

Tertiary Infotech Academy Pte Ltd

UEN: 201200696W

Version 15.0

DOCUMENT VERSION CONTROL RECORD

Version Number	Effective Date of Release	Summary of Included Changes	Author
13.0	Not stated in supplied v13 deck	Legacy SEO for Small and Medium Enterprises reference deck.	Tertiary Infotech Academy Pte Ltd
14.0	16 August 2026	Generative-AI-assisted SEO, prompt frameworks, AI Overviews, multimodal content, agentic governance and ten progressive website labs.	Tertiary Infotech Academy Pte Ltd
15.0	17 August 2026	Added the eight field-tested meta-tag and LinkedIn prompt templates, the five-part prompt skeleton and the pre-publish checklist to the activity prompt packs.	Tertiary Infotech Academy Pte Ltd

TABLE OF CONTENTS

How to Use This Guide.....	4
Shared EcoDesk Live Website.....	4
Learning and outcomes Mapping.....	4
Prompt Frameworks.....	5
Topic 1 — Fundamentals of SEO and Generative AI.....	5
Lab 1 — Audit the EcoDesk Demo Website.....	8
Lab 2 — Design the SIGNAL AI SEO Strategy.....	10
Topic 2 — AI-Powered Keyword Research and Content Strategy.....	12
Lab 3 — Generate and Validate a Keyword Universe.....	14
Lab 4 — Cluster Keywords and Build Content Briefs.....	17
Topic 3 — Creating SEO-Optimised Content with Generative AI.....	20
Lab 5 — Create a People-First SEO Blog Article.....	23
Lab 6 — Generate Image, Video and LinkedIn Assets.....	26
Topic 4 — AI for On-Page SEO and Content Optimisation.....	30
Lab 7 — Generate and Implement On-Page SEO Meta.....	32
Lab 8 — Optimise an Answer Page for Google AI Overviews.....	36
Topic 5 — SEO Performance, Ethics and Continuous Improvement with AI.....	39
Lab 9 — Analyse Search Console and GA4 Data with AI.....	41
Lab 10 — Build a Governed Agentic SEO Workflow.....	44
Assessment Preparation.....	48
Sources and Further Reading.....	48


How to Use This Guide

Use slides for concepts and discussion. Follow detailed procedures here and in each activity folder. All ten activities progressively optimise one shared fictional EcoDesk website, supported by individual prompt packs, tips, sample articles, data and evidence checklists.

Shared EcoDesk Live Website

Use the same published site throughout the course. It includes an SEO-led commercial pitch, product and article pages, a lead magnet, enquiry funnel and WhatsApp-style assistant.

Live site: <https://tertiarycourses.github.io/C11-Generative-AI-for-Search-Engine-Optimization-SEO/>
<https://tertiarycourses.github.io/C11-Generative-AI-for-SEO/>



Safety boundary: Use public, approved or synthetic data. AI output is a draft or hypothesis until a named human verifies and approves it.

Learning and outcomes Mapping

- LO1: Evaluate SEO and internet marketing strategies of a website.
- LO2: Manage keyword research in alignment with SEO objectives.
- LO3: Provide on-page SEO strategies and recommendations.
- LO4: Manage off-page SEO strategies and performance across multiple channels.
- LO5: Monitor Google Analytics and Search Console reports to evaluate SEO performance.

Prompt Frameworks

Use CONTEXT: Context, Objective, Necessary inputs, Task steps, Evidence, eXclusions and Template.
Use SIGNAL for strategy: Search intent, Information advantage, Generative workflow, Network authority, Accessibility and Learning loop.

Topic 1 — Fundamentals of SEO and Generative AI

Key Concepts

- Search: Earn discoverability by helping people and search systems understand useful content.
- Generative AI: Accelerate research and creation while humans own evidence, intent and approval.
- Strategy: Connect audience needs, business outcomes, content assets, authority and measurement.
- Governance: Set data, fact-check, brand, copyright, disclosure and publishing controls.

SEO is a system of useful signals, not a bag of tricks.

Key point: SEO is a system of useful signals, not a bag of tricks. AI increases execution speed; strategy decides whether that speed creates value or spam.

How modern search finds and serves content

1. Discover
2. Crawl
3. Render
4. Index
5. Retrieve
6. Rank and present

The SERP is now a portfolio

- Organic links: Classic results with titles, snippets and rich features.
- AI Overviews / AI Mode: Generated responses grounded in indexed sources and query fan-out.
- Local, image and video: Intent-specific surfaces with their own eligibility signals.
- Paid and shopping: Commercial inventory distinct from organic ranking.

Traditional SEO and generative-search visibility

Perspective A

- Indexable, useful pages
- Intent and relevance
- Authority and links
- Page experience and technical health

Perspective B

- Clear answer units

- Unique evidence and first-party experience
- Entity consistency
- Extractable text plus helpful images/video

White-hat and black-hat choices

- People-first: Create original value for a defined audience.
- Transparent: Use accurate metadata and visible content.
- Manipulative: Cloaking, doorway pages, link schemes and hidden text.
- Scaled abuse: Many unoriginal AI pages made mainly to manipulate rankings.

AI SEO operating cycle

1. Define outcome
2. Ground the model
3. Generate options
4. Validate evidence
5. Edit and approve
6. Publish, measure, learn

AI SEO strategy: SIGNAL

- S — Search intent: Specify audience, job and decision stage.
- I — Information advantage: Add original examples, expertise and proof.
- G — Generative workflow: Use reusable prompts, inputs and quality gates.
- N — Network authority: Build internal links, earned mentions and credible distribution.
- A — Accessibility: Ensure crawl, render, mobile and media accessibility.
- L — Learning loop: Measure outcomes and refresh from evidence.

Three roles for AI

- Copilot: Human directs each task and approves the output.
- Workflow: Automation passes structured inputs between tools.
- Agent: System plans and acts within a bounded goal, tools and approval policy.

Agentic SEO with control points

1. Goal and scope
2. Research
3. Audit
4. Plan
5. Draft or recommend
6. Human approval
7. Publish

8. Monitor and rollback

Human-in-the-loop risk tiers

- Low: Summarise exports, cluster supplied keywords, flag missing metadata.
- Medium: Draft briefs, internal links, schema and social variants for review.
- High: Publish, change templates, edit robots/canonicals or contact external sites.
- Control: Least privilege, preview, approvals, logs, limits and rollback.

Prompt anatomy: CONTEXT

- C — Context: Business, audience, page, product and constraints.
- O — Objective: One measurable task and target outcome.
- N — Necessary inputs: Verified facts, keyword data, page copy and links.
- T — Task steps: Analysis sequence, checks and decision rules.
- E — Evidence: Citations, uncertainty labels and claim ledger.
- X — eXclusions: Prohibited claims, tactics, private data and assumptions.
- T — Template: Exact output fields, length limits and QA checklist.

Pitfalls and countermeasures

- Hallucination: Ground claims and require a fact/uncertainty table.
- Generic content: Add first-party examples and expert review.
- Keyword stuffing: Optimise for intent and natural coverage, not repetition.
- Privacy leakage: Use public or synthetic inputs; redact customer data.
- Copyright risk: Do not imitate living creators or reproduce source text.
- Automation drift: Version prompts, log changes and monitor outcomes.

Lab 1 — Audit the EcoDesk Demo Website

LO1 · A6 · K1 · K6 · Evaluate a website's current SEO and internet-marketing strategy

Goal and Deliverable

Establish an evidence-based baseline for a fictional Singapore sustainable-office-products website before any AI-assisted changes.

- Completed baseline audit, five-item priority backlog and preserved original website

Scenario

EcoDesk is a fictional training business. The baseline intentionally contains realistic SEO strengths and weaknesses so learners can inspect code, content and user experience safely.

EcoDesk wants more qualified enquiries from Singapore SMEs but cannot explain which pages, queries or technical issues should be prioritised.

Primary source: <https://developers.google.com/search/docs/essentials>

Assumptions and Guardrails

- All business names, addresses, prices, testimonials and analytics are synthetic.
- No live Search Console or customer data is required.

- The audit must distinguish observed evidence from an AI inference.
- The AI may recommend but must not edit files automatically in this lab.

Questions

1. Which issues affect crawl/index eligibility, relevance, trust or conversion?
2. Which current strengths should be preserved?
3. Which recommendation has the highest impact-to-effort ratio?
4. What evidence is missing before implementation?

Detailed Step-by-Step Procedure

Step 1: Open and inventory

Serve the sample website locally, list every page and asset, and record the page purpose and intended audience.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Inspect source

Check title, description, canonical, robots directive, headings, links, images, scripts and structured data in the HTML.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Review experience

Test navigation, mobile layout, forms, readability, image loading and keyboard access.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Run the audit prompt

Provide only the copied source and audit criteria to the AI. Require observed evidence, severity, confidence and suggested verification.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Challenge the output

Reject invented facts, unsupported ranking claims and advice that conflicts with the visible source.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Prioritise

Score impact, effort, risk and dependency; select the first five actions.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Capture baseline

Save the original files and the completed scorecard before any edits.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Peer review

Exchange audits and reconcile disagreements by pointing to source evidence.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: Every finding cites observable evidence, separates fact from inference, includes a verification method and maps to a business or user outcome.

Lab 2 — Design the SIGNAL AI SEO Strategy

LO1 · A6 · K1 · K6 · Create an organisation-level AI SEO strategy and governed workflow

Goal and Deliverable

Turn the baseline into a 90-day AI-assisted SEO strategy with outcomes, audiences, content priorities, controls and measurement.

- 90-day SIGNAL strategy and governance page

Scenario

Google permits appropriate use of generative AI but warns that scaled, low-value content created mainly to manipulate rankings can violate spam policies.

EcoDesk has one marketer and one subject-matter expert. Management wants faster content production without publishing generic or unsupported environmental claims.

Primary source: <https://developers.google.com/search/docs/fundamentals/using-gen-ai-content>

Assumptions and Guardrails

- The 90-day budget is S\$6,000 and eight staff-days.

- Only public and synthetic data may enter general-purpose AI tools.
- Publishing, outreach and technical directives require named human approval.
- The initial agentic workflow is read-only and recommendation-only.

Questions

1. What outcome should SEO create beyond traffic?
2. Which information advantage can EcoDesk credibly own?
3. Which tasks are suitable for copilot, workflow or agent mode?
4. Which risks require a stop rule or human approval?

Detailed Step-by-Step Procedure

Step 1: Define outcomes

Translate the business goal into audience, conversion and quality outcomes with baselines.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Apply SIGNAL

Complete Search intent, Information advantage, Generative workflow, Network authority, Accessibility and Learning loop.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Select priorities

Choose a small topic cluster, core landing page and supporting media based on the audit.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Design the workflow

Map inputs, prompts, tools, reviewer, publishing gate, log and rollback for each task.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Classify risk

Place each AI action into low, medium or high risk and specify permissions.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Write prompt policy

Define allowed data, sources, fact checking, copyright, disclosure and prohibited claims.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Build roadmap

Sequence foundations, pilot, measurement and expansion across 90 days.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Publish strategy page

Add a concise strategy summary to the sample website and verify navigation and mobile display.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: The strategy links each initiative to an outcome, source of truth, owner, approval gate, metric and rollback or review trigger.

Topic 2 — AI-Powered Keyword Research and Content Strategy

Key Concepts

- Demand: Keywords reveal language and questions, not guaranteed traffic.
- Intent: Classify what the searcher wants to know, compare, do or buy.
- Topic model: Connect entities, subtopics and follow-up questions into useful coverage.
- Prioritisation: Balance relevance, business value, difficulty, evidence and effort.

A keyword list becomes a strategy only after intent, evidence and priority are added.

Key point: A keyword list becomes a strategy only after intent, evidence and priority are added. AI is strong at expansion and clustering; real search data and business judgement remain the source of truth.

Keyword dimensions

- Head vs long-tail: Breadth and volume versus specificity and conversion context.
- Branded vs non-branded: Known demand versus category discovery.
- Local vs national: Place, proximity and service-area context.
- Question vs action: Learning needs versus comparison, booking or purchase.

Grounded AI keyword workflow

1. Seed from offers and customers
2. Collect Search Console / planner data
3. Expand questions and entities
4. Cluster by intent
5. Validate SERPs
6. Prioritise and map pages

Sources of keyword evidence

- First party: Site search, sales calls, support tickets and Search Console queries.
- Google surfaces: Autocomplete, People Also Ask, related searches and Trends.
- Competitive: Ranking pages, content gaps, backlinks and paid-search language.
- Commercial tools: Planner, Ahrefs, Semrush and alternatives; verify definitions and freshness.

What AI can infer and what it cannot know

Perspective A

- Semantic similarity
- Likely intent
- Question variations
- Entity and modifier patterns

Perspective B

- Current volume without data
- Actual difficulty
- Your conversion rate
- Business priority and legal constraints

Search-intent matrix

- Know: Definitions, explanations, guides and evidence.
- Do: Procedures, tools, templates and checklists.
- Compare: Alternatives, criteria, trade-offs and use cases.
- Buy / visit: Offer, trust, price, availability, location and action.

Prompted cluster method

1. Normalise terms
2. Infer intent
3. Group by shared task
4. Name parent topic
5. Choose pillar or support page

6. Flag cannibalisation

Keyword opportunity score

- Relevance: Does the query fit the offer and audience?
- Evidence: Can we add credible, differentiated value?
- Demand: Is there observable interest or strategic importance?
- Competition: Can the site plausibly earn visibility?
- Conversion: Does the intent support a valued action?
- Effort: What content, technical and authority work is needed?

Content architecture

- Pillar: Broad decision guide connecting the topic cluster.
- Cluster: Focused pages resolving sub-intents without duplication.
- Conversion: Service, product or lead page connected by relevant internal links.

Cannibalisation warning signs

- Same intent: Multiple pages promise the same answer.
- SERP switching: Different URLs rank for one query over time.
- Split authority: Links and engagement fragment across near-duplicates.
- Fix: Consolidate, differentiate purpose, redirect or adjust internal links.

Prompt controls for keyword research

- Input table: Require keyword, source, metric date and landing page.
- No invented metrics: Use null/unknown when volume or difficulty is missing.
- Reason codes: Explain intent and cluster assignment briefly.
- Review queue: Flag ambiguous, sensitive or low-confidence terms.

From keyword to editorial brief

1. Primary intent
2. Audience problem
3. Unique angle
4. Evidence plan
5. Outline and media
6. Internal links
7. Conversion and measurement

Lab 3 — Generate and Validate a Keyword Universe

LO2 · A1 · K2 · K8 · Use AI to expand and validate search terms without inventing metrics

Goal and Deliverable

Create a grounded keyword universe for EcoDesk from supplied first-party notes and synthetic planner/Search Console exports.

- Validated keyword universe and topic-hub webpage

Scenario

AI can expand, normalise and group language, but search volume, competition and current SERP evidence must come from supplied or verified data.

EcoDesk needs a content plan spanning recycled stationery, ergonomic accessories and sustainable workplace practices.

Primary source: <https://support.google.com/google-ads/answer/7337243>

Assumptions and Guardrails

- Keyword metrics are synthetic and dated 16 August 2026.
- Missing volume and difficulty must remain unknown, never estimated by the model.
- Search intent requires manual SERP validation before publishing.
- Singapore is the primary market and English is the primary language.

Questions

1. Which seeds come from real customer language?
2. Which expansions are relevant but unsupported by demand data?
3. Which queries reveal local, comparison or purchase intent?
4. Which terms should be excluded for weak fit or compliance risk?

Detailed Step-by-Step Procedure

Step 1: Review inputs

Inspect the synthetic Search Console, planner and customer-language rows; record source and date.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Create seed sets

Group seeds by product, problem, audience, location and desired outcome.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Run expansion prompt

Require the AI to return phrase, inferred intent, rationale and source seed without metrics.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Normalise

Remove duplicates, spelling noise and meaningless variations while preserving useful language differences.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Join metrics

Merge only the supplied metrics by exact keyword; leave unmatched values blank.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Validate samples

Manually inspect representative SERPs and record observed result types and competing intent.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Filter

Remove irrelevant, misleading, unsupported or very low-value terms.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Build hub

Add the selected themes and questions to keyword-hub.html as crawler-readable static content.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: No metric is AI-invented; every retained keyword has a source, intent, cluster candidate, validation status and business-fit rationale.

Field-Tested Prompt Templates

Bulk generation exposes the cannibalisation risk this lab's keyword universe must prevent. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Bulk metas for a whole section of the site

Where AI genuinely changes the economics: a hundred course or product pages that would take a week by hand take an afternoon with review.

```
You are an SEO assistant. I will paste a table of URLs with columns:
URL, page type, primary keyword, current meta title.
```

```
For each row write a new meta title (max 60 chars) and meta
description (max 155 chars).
```

Rules:

- Keep each row's primary keyword in the first half of the title
- Never reuse the same opener twice - vary the phrasing across rows
- Append " | [BRAND]" only where it still fits the limit
- Flag any two rows that target the same query (cannibalisation risk)
- Return one table I can paste into a spreadsheet

```
[PASTE YOUR TABLE]
```

What to check: Spot-check ten rows, not two. Bulk runs fail in patterns — if the model drifts into a repeated formula halfway down, regenerate from that row with "vary the openers" restated.

In this lab: Paste the EcoDesk page inventory (index, products, guide, article, thank-you). The cannibalisation flag is the graded output: two EcoDesk pages chasing one query is a clustering defect, not a copy defect.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.
- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Lab 4 — Cluster Keywords and Build Content Briefs

LO2 · A1 · K2 · K8 · Map validated keywords to intent-led pages and editorial briefs

Goal and Deliverable

Transform the validated keyword universe into a non-cannibalising topic cluster and three grounded content briefs.

- Topic-cluster architecture and three approved content briefs

Scenario

Useful information architecture groups content by audience tasks and makes important pages reachable through descriptive internal links.

EcoDesk can publish only one pillar page and two supporting articles this quarter, so duplication and weak business fit are expensive.

Primary source: <https://developers.google.com/search/docs/fundamentals/seo-starter-guide>

Assumptions and Guardrails

- The approved cluster is sustainable office setup for Singapore SMEs.
- EcoDesk has access to an operations specialist for first-party examples.
- Environmental claims require documentary support and cautious wording.
- Existing pages may be consolidated when they serve the same intent.

Questions

1. Which terms share the same dominant intent?
2. Which page should be pillar, support or conversion content?
3. Where could two proposed pages cannibalise each other?
4. What unique evidence makes each brief publishable?

Detailed Step-by-Step Procedure

Step 1: Prepare input

Use only validated rows from Lab 3 and add current-page mapping where relevant.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Cluster with AI

Ask for group, parent topic, intent, confidence and ambiguous cases; prohibit invented metrics.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Review ambiguity

Split mixed-intent groups and merge semantic variations that lead to the same page purpose.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Design architecture

Choose pillar, supporting articles and conversion destination with descriptive internal links.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Check cannibalisation

Compare title, intent, promise and target query across every proposed page.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Create briefs

For three pages define audience, answer, angle, evidence, outline, media, links and conversion.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Add QA gates

Include fact, originality, brand, environmental-claim and SEO checks.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Publish plan page

Represent the cluster as static HTML cards with clear navigation and no fake performance claims.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: Every keyword maps to one primary page, each brief has a distinct intent and information advantage, and internal links form a logical user journey.

Field-Tested Prompt Templates

Cluster-to-page mapping is only real once each page has its own non-competing title. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Meta title and description for a single page

The workhorse. Asking for three options in a table forces variety and makes the character counts checkable at a glance.

You are an SEO copywriter. Write 3 meta title options (max 60 characters each) and 3 meta description options (max 155 characters each) for the page below.

Page topic: [WHAT THE PAGE IS ABOUT]

Primary keyword: [KEYWORD] - it must appear near the start of each title

Audience: [WHO SEARCHES FOR THIS]

Differentiator: [WHY THIS PAGE - e.g. price, range, location]

Rules:

- No clickbait and no claims the page cannot back up
- Include one concrete number or specific detail in each description
- End each description with a soft call to action
- Return a table: option, title, title characters, description, description characters

What to check: The model counts characters unreliably — paste the winner into a pixel-width checker before shipping, and confirm every claim is actually true on the page.

In this lab: Run it on the EcoDesk product page. Every claim must trace to visible page content; EcoDesk has no published price, rating, certification or environmental result to cite.

Template 2: Bulk metas for a whole section of the site

Where AI genuinely changes the economics: a hundred course or product pages that would take a week by hand take an afternoon with review.

You are an SEO assistant. I will paste a table of URLs with columns: URL, page type, primary keyword, current meta title.

For each row write a new meta title (max 60 chars) and meta description (max 155 chars).

Rules:

- Keep each row's primary keyword in the first half of the title
- Never reuse the same opener twice - vary the phrasing across rows
- Append " | [BRAND]" only where it still fits the limit
- Flag any two rows that target the same query (cannibalisation risk)
- Return one table I can paste into a spreadsheet

[PASTE YOUR TABLE]

What to check: Spot-check ten rows, not two. Bulk runs fail in patterns — if the model drifts into a repeated formula halfway down, regenerate from that row with "vary the openers" restated.

In this lab: Paste the EcoDesk page inventory (index, products, guide, article, thank-you). The cannibalisation flag is the graded output: two EcoDesk pages chasing one query is a clustering defect, not a copy defect.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.
- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Topic 3 — Creating SEO-Optimised Content with Generative AI

Key Concepts

- Brief: A grounded contract for audience, intent, evidence, angle and outcome.
- Draft: AI proposes structure and language; experts add truth and experience.
- Media: Images and video answer visual intent and improve comprehension.
- Distribution: Social content earns attention and relationships; backlinks must be editorially deserved.

The publishable unit is not an AI draft; it is a verified, edited asset with a reason to exist.

Key point: The publishable unit is not an AI draft; it is a verified, edited asset with a reason to exist.

Original evidence, useful structure and human accountability separate content from commodity output.

Editorial workflow

1. Research and brief
2. Evidence ledger
3. Outline
4. Draft sections
5. Expert edit
6. SEO/media QA
7. Approve and publish

A strong content brief

- Audience: Role, context, knowledge and decision.
- Intent: Primary question and follow-up needs.
- Angle: Specific promise and information advantage.
- Evidence: Claims, sources, examples and subject-matter input.
- Structure: Answer-first headings, comparison, process and FAQ.
- Action: Next useful step and measurement event.

Answer-first does and don'ts

Perspective A

- Lead with a concise answer
- Use descriptive headings
- Show steps, tables and examples
- Cite facts near the claim

Perspective B

- Do not pad the introduction
- Do not repeat the keyword unnaturally
- Do not hide uncertainty
- Do not promise rankings or AI citations

E-E-A-T as an editorial review lens

- Experience: Demonstrate first-hand process, examples and outcomes.
- Expertise: Use qualified authors and technically accurate explanations.
- Authoritativeness: Earn recognition, references and consistent entity information.
- Trust: Make sources, authorship, updates, policies and limitations clear.

Fact-safe drafting loop

1. List claims
2. Attach source or owner
3. Draft with citation placeholders
4. Verify wording
5. Remove unsupported claims
6. Approve claim ledger

Image-generation prompt controls

- Purpose: Define query, page section and learning/conversion job.
- Composition: Specify subject, framing and responsive crop zones.
- Accuracy: List must-have details and prohibited artefacts.
- Rights: Avoid logos, copyrighted characters and creator imitation.
- Delivery: File type, dimensions, compression and descriptive filename.
- Accessibility: Write alt text for function and context, not keyword stuffing.

Image SEO

- Relevant original visual: Illustrate the page's actual topic or evidence.
- Efficient delivery: Responsive sizes, WebP/AVIF where appropriate, dimensions and lazy loading.
- Discoverable context: Descriptive filename, nearby text and image sitemap where needed.
- Useful alt text: Concisely communicate purpose to users who cannot see the image.

Video SEO

- Searchable page: Give each important video a stable, indexable watch page.
- Transcript: Expose accurate spoken content as useful text.
- Metadata: Descriptive title, thumbnail, duration, upload date and VideoObject where valid.
- Chapters: Help users navigate moments and clarify subtopics.

One article to a social campaign

1. Extract evidence and insight
2. Choose audience tension
3. Draft platform-native post
4. Add proof and point of view
5. Invite discussion
6. Disclose destination
7. Track qualified visits

Earned links and manipulative links

Perspective A

- Original research and tools
- Expert commentary
- Useful resources
- Relevant outreach and digital PR

Perspective B

- Bought or exchanged ranking links
- Mass low-quality guest posts
- Automated comment/forum spam
- Invented endorsements or fake profiles

AI content QA scorecard

- Intent: Resolves the real task completely.
- Accuracy: Every material claim is verified.
- Originality: Adds examples, data or judgement unavailable elsewhere.
- Clarity: Scannable, specific and free of filler.
- Brand: Voice, risk and disclosure requirements met.
- SEO: Metadata, links, media and structured data match the visible page.

Lab 5 — Create a People-First SEO Blog Article

LO3 · A5 · K7 · K10 · Generate and edit a grounded long-form article

Goal and Deliverable

Use a staged prompt workflow to create an answer-first article from a verified brief and claim ledger.

- Verified, edited SEO blog article with claim ledger

Scenario

Helpful content is created primarily to benefit people; AI-assisted content still needs originality, accuracy, transparent authorship and appropriate expertise.

EcoDesk wants a practical guide for Singapore SME office managers who need credible sustainability actions without exaggerated claims.

Primary source: <https://developers.google.com/search/docs/fundamentals/creating-helpful-content>

Assumptions and Guardrails

- All examples are illustrative and must be labelled when not source-backed.
- The article must not claim certification, guaranteed savings or environmental impact without evidence.
- The author is an EcoDesk workplace sustainability adviser with a transparent bio.
- Target length is a guideline; completeness and clarity win.

Questions

1. What is the reader's main decision?
2. Which claims need sources or qualification?
3. What first-party experience can make the article non-commodity?
4. Which internal links continue the user journey?

Detailed Step-by-Step Procedure

Step 1: Lock the brief

Confirm audience, intent, angle, evidence and conversion before drafting.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Build claim ledger

List every factual or quantitative claim with source, owner and approved wording.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Generate outline

Ask for answer-first sections, comparison, checklist, examples and FAQ; remove filler headings.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Draft in sections

Generate one section at a time using only approved evidence and visible uncertainty markers.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Expert edit

Add specific operational examples, remove generic phrasing and test whether advice is actionable.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: SEO edit

Refine title, description, headings, internal links, image plan and concise answer blocks.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Run adversarial QA

Ask a separate review prompt to find unsupported claims, contradictions, duplication and manipulation risk.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Publish and inspect

Create the static blog page, validate links and read it on mobile.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: Every material claim is supported or qualified, the page resolves intent, adds original operational value and passes the six-part AI content QA scorecard.

Field-Tested Prompt Templates

Distribution copy is drafted from the finished article, using only claims the article already supports. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or

labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Turn an article into a LinkedIn post

The most common repurposing job. The two rules that matter most: the hook must survive the "...see more" fold, and the post must contain one concrete number so it reads like experience rather than content marketing.

Turn the article below into a LinkedIn post for my company page.

Structure:

- Hook: the first 2 lines must work before the "...see more" fold - no throat-clearing
- Body: 3-5 short paragraphs, one idea each, plain English
- Include one concrete example or number from the article
- CTA: one line inviting readers to the full article at [URL]
- 3-5 hashtags: one broad, the rest niche

Rules: no emojis in the first line, never open with "I'm excited to share", 150-220 words total.

[PASTE ARTICLE OR KEY POINTS]

What to check: The concrete number must come from the article, not from the model. If the article has no number, the honest fix is to add one to the article first.

In this lab: Feed it the EcoDesk guide page. The post must seek earned attention only — no reciprocal-link request, no paid placement, no engagement pod.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.
- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Lab 6 — Generate Image, Video and LinkedIn Assets

LO4 · A3 · A5 · K3 · K5 · Create multimodal SEO assets and an ethical off-page distribution plan

Goal and Deliverable

Prompt text-to-image and video concepts, write accessible metadata, and repurpose the blog into a LinkedIn post designed to earn attention and editorial backlinks.

- Multimodal SEO asset pack and ethical LinkedIn distribution post

Scenario

Images and video should be relevant, accessible and technically efficient. Social distribution can earn discovery and links, but manipulative link schemes remain prohibited.

EcoDesk will launch the new guide with one original hero visual, a 45-second explainer and a LinkedIn post for SME operations leaders.

Primary source: <https://developers.google.com/search/docs/appearance/google-images>

Assumptions and Guardrails

- Generated visuals may not include logos, copyrighted characters or unsupported certifications.
- The video is a concept storyboard; learners do not need paid generation tools.
- The LinkedIn post must not ask for reciprocal or paid links.
- All media files use synthetic examples and are clearly labelled.

Questions

1. What job does each media asset perform?
2. What visual inaccuracies or rights risks must the prompt prevent?
3. What belongs in alt text versus a caption?
4. Why would another professional voluntarily reference the article?

Detailed Step-by-Step Procedure

Step 1: Define media intent

Map one visual and one video to specific article sections and search/user needs.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Write image prompt

Specify subject, composition, accuracy constraints, prohibited elements, crop zones and output use.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Review concept

Check authenticity, text artefacts, representation, copyright and environmental-claim risk.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Prepare image SEO

Write descriptive filename, dimensions, compression plan, alt text and surrounding caption.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Storyboard video

Create six scenes with narration, on-screen text, b-roll and evidence references.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Create video SEO

Draft title, description, transcript, thumbnail brief, chapters and valid VideoObject fields.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Draft LinkedIn post

Run Templates 1-4 in the prompt pack: repurpose the article, write a first-person post, adapt the announcement to the guide launch and plan a five-post week. Lead with a practitioner insight, cite one useful example, link transparently and invite informed discussion.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Publish library

Add static media cards, transcript and metadata examples to the sample site.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: Assets are relevant, accessible, rights-safe and evidence-aligned; the LinkedIn post offers standalone value and seeks earned rather than manipulated links.

Field-Tested Prompt Templates

Off-page reach is earned with standalone value; these four cover the post types this lab produces. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Turn an article into a LinkedIn post

The most common repurposing job. The two rules that matter most: the hook must survive the "...see more" fold, and the post must contain one concrete number so it reads like experience rather than content marketing.

Turn the article below into a LinkedIn post for my company page.

Structure:

- Hook: the first 2 lines must work before the "...see more" fold - no throat-clearing
- Body: 3-5 short paragraphs, one idea each, plain English
- Include one concrete example or number from the article
- CTA: one line inviting readers to the full article at [URL]
- 3-5 hashtags: one broad, the rest niche

Rules: no emojis in the first line, never open with "I'm excited to share", 150-220 words total.

[PASTE ARTICLE OR KEY POINTS]

What to check: The concrete number must come from the article, not from the model. If the article has no number, the honest fix is to add one to the article first.

In this lab: Feed it the EcoDesk guide page. The post must seek earned attention only — no reciprocal-link request, no paid placement, no engagement pod.

Template 2: A first-person post from something you learned

Personal posts outperform company posts on LinkedIn, and the guardrail line here is the difference between authentic and invented.

Write a first-person LinkedIn post based on this experience:

What happened: [1-3 SENTENCES ON WHAT YOU DID OR LEARNED]

Who it helps: [AUDIENCE]

The takeaway: [THE ONE THING READERS SHOULD REMEMBER]

Rules:

- Open with the moment of surprise or failure, not the lesson
- Short lines, generous white space, max 3 hashtags
- Do NOT invent details - if a fact is missing, ask me instead
- End with a genuine question that invites comments, not "thoughts?"
- 120-180 words

What to check: "Do NOT invent details — ask me instead" is the line to test. If the model fills a gap silently instead of asking, the prompt failed and the output is unusable.

In this lab: Use your own real experience from this lab as the input. A first-person post attributed to a person must not be built from fictional EcoDesk events.

Template 3: Announcing a course run, webinar or event

Announcement posts die when they lead with the thing instead of the problem. This template forces the order right and forces the order right.

Write a LinkedIn post announcing this course run for a company page.

Course: [COURSE NAME]
Date and format: [DATE, CLASSROOM / ONLINE]
Who it is for: [AUDIENCE]
Registration link: [URL]

Rules:

- Lead with the problem the course solves, not the course name
- One line on what participants build or walk away with
- Single CTA line with the link, under 150 words, max 4 hashtags

What to check: Every concrete claim in the post must be verifiable against the live page before posting.

In this lab: Adapt it to the EcoDesk guide launch instead of a course. Any figure you cannot evidence from the page is an invented claim.

Template 4: One article, five posts — a week of content

Repurposing is where consistency comes from. Asking the model to vary the format across the week stops your feed reading like a template.

Take the article below and plan 5 LinkedIn posts from it, one per weekday.

For each post give: the angle, the opening hook line, and 2-3 bullet points of body content. Vary the format across the week: one how-to, one contrarian take, one mini case study, one list post, one question post. Never repeat the same hook structure twice.

Flag which single post you would put paid spend behind, and why.

[PASTE ARTICLE]

What to check: The paid-spend recommendation is a judgement, not a finding — the model has no performance data. Treat it as one hypothesis to test against your own past post metrics.

In this lab: The "mini case study" slot is the trap: EcoDesk has no customer results, so that post must be reframed as a worked example or dropped with a reason recorded.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.
- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Topic 4 — AI for On-Page SEO and Content Optimisation

Key Concepts

- Page signal: Align title, heading, copy, media, links and structured data with one useful purpose.
- Technical eligibility: Search must be allowed to crawl, render, index and show a snippet.
- Information architecture: Make priority pages reachable and relationships explicit.
- AI features: Foundational SEO and unique, people-first content remain the core.

There is no magic AI Overview schema.

Key point: There is no magic AI Overview schema. A page must first be indexable and eligible for Google Search; clear, unique and reliable content does the rest.

On-page signal stack

- Snippet: Title and description set expectation; Google may rewrite them.
- Structure: One clear H1 with descriptive H2/H3 sections.
- Content: Complete, accurate answer with original evidence.
- Media: Helpful images/video with accessible context.
- Links: Descriptive internal links and cited external evidence.
- Markup: Structured data that matches visible content.

Metadata prompt workflow

1. Supply page facts and query
2. Draft 5 distinct options
3. Check intent and uniqueness
4. Enforce length as guidance
5. Remove unsupported claims
6. Human select and edit

Title and description quality

- Specific: Name the page's value and distinguishing context.
- Accurate: Do not promise content or offers absent from the page.
- Unique: Avoid sitewide boilerplate and near duplicates.
- Readable: Natural language first; no keyword lists or excessive punctuation.

URL, canonical and index controls

- URL: Stable, concise, descriptive and consistent.
- Canonical: Indicate preferred duplicate while keeping internal signals aligned.
- Noindex: Remove a page from search; do not rely on robots.txt for this.
- Robots: Control crawling, not guaranteed indexing removal.

Site architecture

1. Home

2. Topic hub
3. Category / pillar
4. Supporting article
5. Conversion page
6. Related resources

Internal-link decisions

- Relevance: Link where the destination helps the reader continue.
- Anchor: Describe the destination naturally.
- Priority: Make valuable pages reachable in few logical steps.
- Maintenance: Repair broken, redirected and orphaned links.

Structured data rules

- Eligible type: Use a Google-supported type only when the page qualifies.
- Visible match: Markup must reflect content users can see.
- Required properties: Include valid, current values and stable URLs.
- Validation: Test syntax and monitor enhancements; rich results are not guaranteed.

AI Overview / AI Mode eligibility

1. Allow crawl and render
2. Index with snippet eligibility
3. Resolve intent with unique value
4. Use clear textual answer units
5. Support with media and evidence
6. Earn selection; no guarantee

Query fan-out implications

- Subtopics: Cover the natural supporting questions without padding.
- Entities: Use consistent names and relationships.
- Comparisons: Provide criteria, trade-offs and decision context.
- Freshness: Update material facts and show meaningful review dates.

Useful GEO/AEO practice and unsupported hacks

Perspective A

- Foundational technical SEO
- People-first, non-commodity content
- Clear structure and first-party evidence
- Search Console measurement

Perspective B

- Special AI-only markup
- Guaranteed citation formulas
- Inauthentic mentions
- Publishing an llms.txt file as a Google ranking tactic

Core Web and mobile hygiene

- LCP: Deliver the main content quickly with efficient media.
- INP: Keep interaction responsive by controlling script work.
- CLS: Reserve space for images, embeds and dynamic UI.
- Mobile: Readable, usable and equivalent primary content across devices.

Lab 7 — Generate and Implement On-Page SEO Meta

LO3 · A2 · K10 · Use prompts to create and implement accurate on-page elements

Goal and Deliverable

Generate metadata options from visible page facts, then implement HTML, internal links, image attributes and valid structured data.

- Optimised static product page and metadata decision log

Scenario

Titles and snippets should accurately represent visible content. Structured data must match the page and does not guarantee a rich result.

The EcoDesk product page is useful but generic. The team needs stronger search clarity without inventing prices, ratings, stock or environmental claims.

Primary source: <https://developers.google.com/search/docs/appearance/title-link>

Assumptions and Guardrails

- Product names and descriptions are fictional.
- No reviews, aggregate rating or certification markup may be added.
- The canonical URL is a training placeholder and must be changed before real deployment.
- Title and description lengths are editorial guidance, not ranking rules.

Questions

1. Which page facts can safely appear in metadata?
2. Which generated option best matches intent and brand?
3. Which structured-data type is eligible and visible?
4. Which internal links reduce user effort?

Detailed Step-by-Step Procedure

Step 1: Extract page facts

Create a fact table from visible product content and approved business information.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Generate options

Run Template 1 (single page) from the prompt pack for five titles and descriptions with rationale, distinct angle and prohibited-claim check, then Template 3 (local intent) to test whether a Singapore or trust signal is honestly available.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Select and edit

Choose one option, improve natural language and confirm it matches the page.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Implement head elements

Add title, description, canonical, robots, Open Graph and social-card metadata, verifying each length in a pixel-width checker rather than trusting the model count.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Improve structure

Use one clear H1, descriptive sections, useful copy and accessible image attributes.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Add internal links

Connect the page to the pillar, blog and enquiry path with descriptive anchors.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Add JSON-LD

Use eligible visible product/organisation facts only and keep JSON valid.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Validate

Inspect source, HTML, mobile layout and structured data; record limitations.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: Metadata is unique and accurate, visible content and JSON-LD agree, links work, images are accessible and no unsupported rating, price or guarantee is present.

Field-Tested Prompt Templates

On-page metadata is this lab's core deliverable, at single-page, section and local-intent scale. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Meta title and description for a single page

The workhorse. Asking for three options in a table forces variety and makes the character counts checkable at a glance.

```
You are an SEO copywriter. Write 3 meta title options (max 60
characters each) and 3 meta description options (max 155 characters
each) for the page below.
```

```
Page topic: [WHAT THE PAGE IS ABOUT]
Primary keyword: [KEYWORD] - it must appear near the start of each title
Audience: [WHO SEARCHES FOR THIS]
Differentiator: [WHY THIS PAGE - e.g. price, range, location]
```

Rules:

- No clickbait and no claims the page cannot back up
- Include one concrete number or specific detail in each description
- End each description with a soft call to action
- Return a table: option, title, title characters, description, description characters

What to check: The model counts characters unreliably — paste the winner into a pixel-width checker before shipping, and confirm every claim is actually true on the page.

In this lab: Run it on the EcoDesk product page. Every claim must trace to visible page content;

EcoDesk has no published price, rating, certification or environmental result to cite.

Template 2: Bulk metas for a whole section of the site

Where AI genuinely changes the economics: a hundred course or product pages that would take a week by hand take an afternoon with review.

You are an SEO assistant. I will paste a table of URLs with columns:
URL, page type, primary keyword, current meta title.

For each row write a new meta title (max 60 chars) and meta description (max 155 chars).

Rules:

- Keep each row's primary keyword in the first half of the title
- Never reuse the same opener twice - vary the phrasing across rows
- Append " | [BRAND]" only where it still fits the limit
- Flag any two rows that target the same query (cannibalisation risk)
- Return one table I can paste into a spreadsheet

[PASTE YOUR TABLE]

What to check: Spot-check ten rows, not two. Bulk runs fail in patterns — if the model drifts into a repeated formula halfway down, regenerate from that row with "vary the openers" restated.

In this lab: Paste the EcoDesk page inventory (index, products, guide, article, thank-you). The cannibalisation flag is the graded output: two EcoDesk pages chasing one query is a clustering defect, not a copy defect.

Template 3: Local-intent metas for Singapore searches

Singapore searchers respond to local trust signals the way other markets respond to discounts — but only claim them where they are true.

Write meta titles and descriptions for this page targeting searchers in Singapore.

Page: [PAGE TOPIC]

Primary keyword: [KEYWORD] singapore

Rules:

- Mention Singapore naturally in the title, not bolted on
- Include one trust signal in the description, e.g. a credential, guarantee or independent rating, ONLY if true for this page
- Max 60 / 155 characters
- Give 3 pairs ranked by expected click-through rate, with one line explaining your ranking

What to check: Trust claims are regulated territory — verify every credential, guarantee and rating against the live page before publishing. A wrong claim in a meta description is a complaint waiting to happen.

In this lab: EcoDesk is a fictional product retailer with no credential or certification of any kind, so the correct learner output is to strike the trust-signal rule and say why. Reinstating it is a quality failure.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.
- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Lab 8 — Optimise an Answer Page for Google AI Overviews

LO3 · A2 · K4 · K10 · Apply foundational SEO and answer-first content for generative search

Goal and Deliverable

Create a clear, source-grounded page that can be indexed and understood by users and generative search features without relying on unsupported GEO hacks.

- Answer-first guide and AI-feature eligibility checklist

Scenario

Google says existing SEO best practices remain relevant for AI features, pages need no special AI markup, and inclusion is never guaranteed.

EcoDesk wants to answer a complex planning query while avoiding inflated promises about appearing in AI Overviews or AI Mode.

Primary source: <https://developers.google.com/search/docs/fundamentals/ai-optimization-guide>

Assumptions and Guardrails

- The page targets Singapore SME office managers.
- The guide uses public general information and synthetic EcoDesk examples.
- FAQ structured data is omitted unless current eligibility and visible-content requirements are met.
- No llms.txt or special AI-only schema is presented as a Google ranking factor.

Questions

1. What direct answer should appear first?
2. Which follow-up subqueries might query fan-out retrieve?
3. What original evidence or experience improves the page?
4. Which technical controls could accidentally remove eligibility?

Detailed Step-by-Step Procedure

Step 1: Define the complex query

State audience, decision and boundaries; list comparison and follow-up needs.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Use query fan-out prompt

Generate related subquestions as hypotheses, then validate relevance manually.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Build answer units

Write concise self-contained answers followed by explanation, table, steps or evidence.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Add information advantage

Include an original workplace checklist and specific Singapore implementation considerations.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Cite and qualify

Place sources near material claims and state limitations clearly.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Check technical eligibility

Confirm crawl, index, snippet, canonical, text availability, internal links and mobile experience.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Reject myths

Document why special AI markup, citation guarantees and inauthentic mentions are excluded.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Publish and test

Open source, check links, headings, visible text and responsive layout.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.

5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: The page is indexable, resolves the main and follow-up intents with unique value, cites evidence and makes no guarantee of AI Overview inclusion.

Field-Tested Prompt Templates

An answer page still needs a snippet that states the answer honestly before the click. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Meta title and description for a single page

The workhorse. Asking for three options in a table forces variety and makes the character counts checkable at a glance.

```
You are an SEO copywriter. Write 3 meta title options (max 60
characters each) and 3 meta description options (max 155 characters
each) for the page below.
```

```
Page topic: [WHAT THE PAGE IS ABOUT]
Primary keyword: [KEYWORD] - it must appear near the start of each title
Audience: [WHO SEARCHES FOR THIS]
Differentiator: [WHY THIS PAGE - e.g. price, range, location]
```

Rules:

- No clickbait and no claims the page cannot back up
- Include one concrete number or specific detail in each description
- End each description with a soft call to action
- Return a table: option, title, title characters, description, description characters

What to check: The model counts characters unreliably — paste the winner into a pixel-width checker before shipping, and confirm every claim is actually true on the page.

In this lab: Run it on the EcoDesk product page. Every claim must trace to visible page content; EcoDesk has no published price, rating, certification or environmental result to cite.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.

- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Topic 5 — SEO Performance, Ethics and Continuous Improvement with AI

Key Concepts

- Visibility: Search impressions and cited presence show discoverability, not business value alone.
- Engagement: Qualified visits and useful actions indicate whether intent was met.
- Conversion: Connect organic journeys to leads, sales or other valued outcomes.
- Learning: Use baselines, hypotheses and controlled changes before attributing impact.

AI can explain a dashboard; it cannot repair weak measurement.

Key point: AI can explain a dashboard; it cannot repair weak measurement. Clean definitions, complete inputs and cautious attribution come before generated insights.

Measurement stack

- Search Console: Queries, pages, clicks, impressions, CTR, position and indexing.
- GA4: Sessions, users, engagement, events, key events and landing-page outcomes.
- Technical crawl: Status, metadata, canonicals, links, directives and performance.
- Business system: Lead quality, revenue, margin, retention and offline outcome.

Search Console and GA4 answer different questions

Perspective A

- Did Google show and send clicks?
- Which query/page combinations changed?
- Is the URL indexed and eligible?
- Which search features are present?

Perspective B

- What did visitors do after arrival?
- Which landing pages drove key events?
- How do channels and campaigns compare?
- What is the user/session journey?

SEO KPI tree

- Reach: Impressions and share of eligible visibility.
- Visit: Clicks, organic sessions and click-through rate.
- Experience: Engagement, scroll, content use and page performance.
- Outcome: Key events, leads, revenue or qualified actions.
- Efficiency: Cost, production time and refresh throughput.
- Quality guardrails: Accuracy, complaints, manual actions and brand risk.

AI-assisted analysis

1. Export clean data
2. Define metric dictionary
3. Ask one diagnostic question
4. Request calculations and evidence
5. Challenge alternative causes
6. Validate in source
7. Record decision

Avoid attribution traps

- Seasonality: Compare relevant periods and known demand cycles.
- SERP change: AI features, ads and layout can change clicks without ranking loss.
- Tracking change: Consent, tagging and definitions can create artificial shifts.
- Content mix: Averages hide landing-page and intent differences.

Experiment discipline

1. Baseline
2. Hypothesis
3. Target pages and guardrails
4. One coherent change
5. Observation window
6. Compare and interpret
7. Keep, revise or revert

AI visibility measurement

- Official: Use Search Console reports where available and overall performance data.
- Manual: Document query, locale, device, date and cited sources.
- Third-party: Treat estimates as directional and check methodology.
- Business: Measure AI-referred engagement and conversion where identifiable.

Agentic monitoring controls

- Read-only first: Let agents detect and recommend before they can edit.
- Threshold: Require material change and minimum evidence before alerts.
- Approval: Human approves content, technical and outreach actions.
- Audit: Log input, model, prompt, output, reviewer, decision and rollback.

Refresh decisions

- Maintain: Content remains accurate, useful and aligned.

- Improve: Update evidence, intent coverage, media, links or conversion path.
- Consolidate / retire: Merge overlap, redirect obsolete pages or noindex when appropriate.

Responsible AI checklist

- Privacy: No personal, confidential or licensed data without authority.
- Accuracy: Verify claims, calculations and schema against sources.
- Fairness: Check stereotypes, exclusions and inappropriate targeting.
- Transparency: Document AI use where policy or audience trust requires it.
- Ownership: Named reviewer and publisher remain accountable.
- Security: Least privilege, secret handling and safe external actions.

Continuous improvement loop

1. Observe
2. Diagnose
3. Prioritise
4. Prompt and produce
5. Review and publish
6. Measure
7. Learn and standardise

Lab 9 — Analyse Search Console and GA4 Data with AI

LO5 · A3 · A4 · K3 · K4 · K9 · Evaluate SEO performance across channels using defined metrics

Goal and Deliverable

Use synthetic Search Console and GA4 exports to diagnose performance without confusing impressions, clicks, sessions, users or conversions.

- Evidence-backed SEO performance review and experiment plan

Scenario

Search Console describes Google search visibility and clicks; GA4 describes user/session behaviour after arrival. Their scopes and counts need not match.

EcoDesk impressions rose while clicks and enquiries fell. Management wants an explanation and action plan, not a confident AI story.

Primary source: <https://support.google.com/analytics/answer/14731736>

Assumptions and Guardrails

- All metrics are synthetic and contain a planned tracking change in week 5.
- No causal claim may be made from correlation alone.
- AI calculations must be checked against spreadsheet formulas.

- Search features and seasonality are plausible alternative causes.

Questions

1. Which metrics changed materially and at what page/query level?
2. Could definition or tracking changes explain the pattern?
3. Which hypotheses fit all available evidence?
4. What test would distinguish competing explanations?

Detailed Step-by-Step Procedure

Step 1: Define metrics

Create a dictionary for impression, click, CTR, position, user, session, engaged session and key event.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Quality-check data

Inspect date ranges, missing rows, duplicates, aggregation and the week-5 tracking note.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Calculate baseline

Compute weekly and page-level changes with spreadsheet formulas.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Prompt for diagnosis

Provide the dictionary and tables; require calculations, evidence rows, alternatives and confidence. For any high-impression low-CTR page, run the CTR-rescue template and record the diagnosis as a hypothesis with a three-week test window.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Verify

Recalculate samples and reject any unsupported causal statement.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Segment

Compare landing page, query intent, device and channel to find hidden differences.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Design experiment

Choose one coherent change with target pages, guardrails, window and decision rule.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Publish review

Present the KPI tree, evidence, hypotheses and experiment on a static dashboard page.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: Every conclusion traces to supplied data, metric scopes are correct, alternative explanations are considered and the recommended test has a baseline and decision rule.

Field-Tested Prompt Templates

Search Console impressions and CTR are the inputs; the rewrite is the testable next action. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Rescuing a page with high impressions but low CTR

Pull the numbers from Google Search Console first — this prompt turns a reporting insight into a rewrite, which is the single fastest SEO win available to most sites.

Here is Google Search Console data for one page (last 3 months):

Query: [TOP QUERY] Impressions: [N] Clicks: [N]
CTR: [N]% Average position: [N]

Current meta title: [TITLE]
Current meta description: [DESCRIPTION]

The page ranks well but searchers skip it. Diagnose the likely reason in 2 sentences, then write 3 alternative title/description pairs for the same query: one leading with a number, one leading with the

outcome, one leading with a question. Max 60 / 155 characters.

What to check: The diagnosis is a hypothesis, not a fact. Ship one variant, wait three weeks, and compare CTR in Search Console before rolling the pattern out.

In this lab: Use the supplied synthetic Search Console rows only. Record the diagnosis as an inference with a stated confidence level and a three-week measurement window — never as a finding.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.
- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Lab 10 — Build a Governed Agentic SEO Workflow

Goal and Deliverable

Integrate the course into a capstone workflow that monitors, recommends and prepares changes while humans retain control of publishing and outreach.

- Governed agentic SEO capstone and monitoring plan

Scenario

Agentic SEO can coordinate multi-step research, audit and monitoring, but high-risk actions need scoped permissions, approvals, logs and rollback.

EcoDesk wants a weekly agent to inspect synthetic Search Console data and the static site, recommend refreshes and draft pull-request-ready changes without publishing them.

Primary source: <https://developers.google.com/search/docs/essentials/spam-policies>

Assumptions and Guardrails

- The agent has read-only access to analytics and repository content.
- It cannot publish, send outreach, edit robots/canonicals or expose secrets.
- Every recommendation must cite a page, metric row or policy source.
- A human owner approves, edits and implements accepted changes.

Questions

1. What is the bounded goal and stop condition?
2. Which tools and data does the agent actually need?
3. Which actions need preview, approval or prohibition?
4. How will quality, drift, value and incidents be measured?

Detailed Step-by-Step Procedure

Step 1: Set goal and scope

Define the weekly outcome, included pages, excluded actions and measurable success.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 2: Map inputs

List site files, metrics, prompt versions, policy sources and their data classification.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 3: Design stages

Connect research, audit, prioritisation, draft recommendation, validation and approval.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 4: Assign permissions

Use least privilege; mark read, suggest, draft, approve, publish and rollback rights.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 5: Write prompt contract

Require evidence, confidence, policy check, change preview, tests and a no-action option.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 6: Build approval gates

Place human decisions before content, technical, outreach and publication actions.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 7: Create audit log

Record run ID, inputs, model, prompt, findings, reviewer, decision, changes and rollback.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 8: Run tabletop

Simulate hallucinated traffic data and a proposed noindex change; verify the workflow stops safely.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 9: Plan monitoring

Track accepted recommendations, false positives, time saved, outcome change and incidents for 30 days.

Checkpoint: Save the evidence, file or decision produced before continuing.

Step 10: Publish governance page

Present the bounded workflow, permissions and review cadence in the capstone website.

Checkpoint: Save the evidence, file or decision produced before continuing.

Run and Test the Website

1. Keep this activity folder open for its prompts and evidence files.
2. From the repository root, run: `python3 -m http.server 8000 --directory activities/ecodesk-demo-website`
3. Open `http://localhost:8000/` or the published EcoDesk URL.
4. Apply this activity to the same shared website used in all ten activities.
5. Test links, forms, assistant, mobile navigation, keyboard access and page source.
6. Complete the evidence checklist and human approval record.

Test it

Pass condition: The workflow is evidence-bound, read-only by default, blocks high-risk autonomous actions, records decisions and demonstrates a successful stop/rollback tabletop.

Field-Tested Prompt Templates

Repeatable, high-volume jobs are the ones worth governing with approval gates. Source: [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#)

Every template follows the same five-part skeleton: Role — Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.; Context — The page, audience, keyword and differentiator. The model cannot know your business; feed it.; Constraints — Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.; Format — Ask for a table or labelled options so the output drops straight into your workflow.; Guardrail — An explicit instruction not to invent facts, and to ask when information is missing.

Template 1: Bulk metas for a whole section of the site

Where AI genuinely changes the economics: a hundred course or product pages that would take a week by hand take an afternoon with review.

```
You are an SEO assistant. I will paste a table of URLs with columns:
URL, page type, primary keyword, current meta title.
```

```
For each row write a new meta title (max 60 chars) and meta
description (max 155 chars).
```

Rules:

- Keep each row's primary keyword in the first half of the title
- Never reuse the same opener twice - vary the phrasing across rows
- Append " | [BRAND]" only where it still fits the limit
- Flag any two rows that target the same query (cannibalisation risk)

- Return one table I can paste into a spreadsheet

[PASTE YOUR TABLE]

What to check: Spot-check ten rows, not two. Bulk runs fail in patterns — if the model drifts into a repeated formula halfway down, regenerate from that row with "vary the openers" restated.

In this lab: Paste the EcoDesk page inventory (index, products, guide, article, thank-you). The cannibalisation flag is the graded output: two EcoDesk pages chasing one query is a clustering defect, not a copy defect.

Template 2: One article, five posts — a week of content

Repurposing is where consistency comes from. Asking the model to vary the format across the week stops your feed reading like a template.

Take the article below and plan 5 LinkedIn posts from it, one per weekday.

For each post give: the angle, the opening hook line, and 2-3 bullet points of body content. Vary the format across the week: one how-to, one contrarian take, one mini case study, one list post, one question post. Never repeat the same hook structure twice.

Flag which single post you would put paid spend behind, and why.

[PASTE ARTICLE]

What to check: The paid-spend recommendation is a judgement, not a finding — the model has no performance data. Treat it as one hypothesis to test against your own past post metrics.

In this lab: The "mini case study" slot is the trap: EcoDesk has no customer results, so that post must be reframed as a worked example or dropped with a reason recorded.

Pre-Publish Checklist

- Character/pixel limits verified in a real checker — Models miscount; Google truncates at pixel width, not characters.
- Every number and claim confirmed true — Fabricated specifics are the most common AI failure — and the most damaging.
- Primary keyword present and early — The one on-page signal a meta rewrite must never lose.
- Reads differently from your last five outputs — Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
- A named human edited it — Generation is a drafting step, not a publishing pipeline.

Assessment Preparation

- Oral Questioning / assessor clarification: up to 15 minutes where required.
- Open book: course slides, Learner Guide and approved lab materials only.

3. Upload completed answers on the LMS.

Sources and Further Reading

Google: AI features and your website: <https://developers.google.com/search/docs/appearance/ai-features>

Google: Optimizing for generative AI features:
<https://developers.google.com/search/docs/fundamentals/ai-optimization-guide>

Google: Guidance on generative AI content:
<https://developers.google.com/search/docs/fundamentals/using-gen-ai-content>

Google: Spam policies: <https://developers.google.com/search/docs/essentials/spam-policies>

Google Search Essentials: <https://developers.google.com/search/docs/essentials>

Google structured data guidelines: <https://developers.google.com/search/docs/appearance/structured-data/sd-policies>

Google Search Console performance: <https://support.google.com/webmasters/answer/7576553>

Google Analytics acquisition: <https://support.google.com/analytics/answer/14731736>

Conductor: AI Overview strategies: <https://www.conductor.com/academy/optimization-strategies-google-ai-overviews/>

AllAble: Agentic SEO: <https://www.allable.ai/blog/what-is-agentic-seo/>

Authoricy: Agentic SEO: <https://www.authoricy.com/blog/agentic-seo>

Opinly: Generative AI for SEO: <https://opinly.ai/blog/using-generative-ai-for-seo>

Gradient Group: SEO prompt examples: <https://gradientgroup.com/28-expert-level-prompt-examples-for-seo/>

Techpresso: SEO prompts: <https://academy.techpresso.co/prompts/chatgpt-prompts-seo>

User-supplied source register: <https://www.tertiarycourses.com.sg/generative-ai-for-search-engine-optimization-seo.html>
<https://www.tertiarycourses.com.sg/generative-ai-for-seo.html>