

Prompt Pack - Lab 3: Generate and Validate a Keyword Universe

Replace bracketed fields with verified inputs. Never paste personal, confidential, licensed or credential data into a tool without authority.

Prompt 1 - CONTEXT task prompt

Context: You support EcoDesk, a fictional Singapore sustainable-office-products business. Audience: [role]. Page/task: [page]. Brand facts: [verified facts only].

Objective: [one measurable outcome].

Necessary inputs: [paste source text/data with source and date].

Task steps: Expand the supplied seed keywords and cluster by intent. Never invent volume, difficulty, CPC or trend data; use unknown where input is absent.

Evidence: Cite the exact supplied row, page element or source for every material conclusion. Separate observation, inference and recommendation.

eXclusions: No invented metrics, rankings, citations, testimonials, prices, certifications, environmental benefits or legal claims. No keyword stuffing, link schemes or private data.

Template: Return findings/options, evidence, uncertainty, recommendation, human checks and acceptance checklist.

Prompt 2 - Adversarial quality review

Act as a critical SEO editor. Review intent, accuracy, originality, clarity, brand, accessibility, Google Search policies and the lab acceptance test. Return issue, severity, evidence, correction and reviewer decision.

Prompt 3 - Human approval summary

Prepare a change preview. Show affected files/fields, before/after meaning, evidence, limitations, tests, risks, rollback and an approve/revise/reject recommendation. Do not publish.

Field-tested prompt templates

Published templates from [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#). Bulk generation exposes the cannibalisation risk this lab's keyword universe must prevent.

Every template follows the same five-part skeleton - it is the skeleton, not the tool, that does the work:

Part	What it does
Role	Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.
Context	The page, audience, keyword and differentiator. The model cannot know your business; feed it.
Constraints	Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.
Format	Ask for a table or labelled options so the output drops straight into your workflow.
Guardrail	An explicit instruction not to invent facts, and to ask when information is missing.

Replace the [BRACKETED] fields with verified inputs before running any template below.

Template 1 - Bulk metas for a whole section of the site

Where AI genuinely changes the economics: a hundred course or product pages that would take a week by hand take an afternoon with review.

You are an SEO assistant. I will paste a table of URLs with columns:
URL, page type, primary keyword, current meta title.

For each row write a new meta title (max 60 chars) and meta description (max 155 chars).

Rules:

- Keep each row's primary keyword in the first half of the title
- Never reuse the same opener twice - vary the phrasing across rows
- Append " | [BRAND]" only where it still fits the limit
- Flag any two rows that target the same query (cannibalisation risk)
- Return one table I can paste into a spreadsheet

[PASTE YOUR TABLE]

What to check: Spot-check ten rows, not two. Bulk runs fail in patterns - if the model drifts into a repeated formula halfway down, regenerate from that row with "vary the openers" restated.

In this lab: Paste the EcoDesk page inventory (index, products, guide, article, thank-you). The cannibalisation flag is the graded output: two EcoDesk pages chasing one query is a clustering defect, not a copy defect.

Pre-publish checklist

Whatever the prompt produced, it goes through the same 60-second review before it ships.

Check	Why it matters
Character/pixel limits verified in a real checker	Models miscount; Google truncates at pixel width, not characters.
Every number and claim confirmed true	Fabricated specifics are the most common AI failure - and the most damaging.
Primary keyword present and early	The one on-page signal a meta rewrite must never lose.
Reads differently from your last five outputs	Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
A named human edited it	Generation is a drafting step, not a publishing pipeline.

These templates are model-agnostic: the same wording works in ChatGPT, Claude, Gemini and Copilot.

Prompt usage record

Field	Record
Tool/model	
Date/time	
Prompt version	v15.0
Input sources	

Field	Record
Reviewer	
Decision	Approve / Revise / Reject
Human changes	