

# Prompt Pack - Lab 4: Cluster Keywords and Build Content Briefs

Replace bracketed fields with verified inputs. Never paste personal, confidential, licensed or credential data into a tool without authority.

## Prompt 1 - CONTEXT task prompt

Context: You support EcoDesk, a fictional Singapore sustainable-office-products business. Audience: [role]. Page/task: [page]. Brand facts: [verified facts only].

Objective: [one measurable outcome].

Necessary inputs: [paste source text/data with source and date].

Task steps: Cluster only validated keywords into non-overlapping page intents. Produce a pillar/support/conversion map and three grounded editorial briefs.

Evidence: Cite the exact supplied row, page element or source for every material conclusion. Separate observation, inference and recommendation.

eXclusions: No invented metrics, rankings, citations, testimonials, prices, certifications, environmental benefits or legal claims. No keyword stuffing, link schemes or private data.

Template: Return findings/options, evidence, uncertainty, recommendation, human checks and acceptance checklist.

## Prompt 2 - Adversarial quality review

Act as a critical SEO editor. Review intent, accuracy, originality, clarity, brand, accessibility, Google Search policies and the lab acceptance test. Return issue, severity, evidence, correction and reviewer decision.

## Prompt 3 - Human approval summary

Prepare a change preview. Show affected files/fields, before/after meaning, evidence, limitations, tests, risks, rollback and an approve/revise/reject recommendation. Do not publish.

## Field-tested prompt templates

Published templates from [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#). Cluster-to-page mapping is only real once each page has its own non-competing title.

Every template follows the same five-part skeleton - it is the skeleton, not the tool, that does the work:

| Part        | What it does                                                                                         |
|-------------|------------------------------------------------------------------------------------------------------|
| Role        | Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.       |
| Context     | The page, audience, keyword and differentiator. The model cannot know your business; feed it.        |
| Constraints | Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop. |
| Format      | Ask for a table or labelled options so the output drops straight into your workflow.                 |
| Guardrail   | An explicit instruction not to invent facts, and to ask when information is missing.                 |

Replace the [BRACKETED] fields with verified inputs before running any template below.

### Template 1 - Meta title and description for a single page

The workhorse. Asking for three options in a table forces variety and makes the character counts checkable at a glance.

You are an SEO copywriter. Write 3 meta title options (max 60 characters each) and 3 meta description options (max 155 characters each) for the page below.

Page topic: [WHAT THE PAGE IS ABOUT]

Primary keyword: [KEYWORD] - it must appear near the start of each title

Audience: [WHO SEARCHES FOR THIS]

Differentiator: [WHY THIS PAGE - e.g. price, funding, location]

Rules:

- No clickbait and no claims the page cannot back up
- Include one concrete number or specific detail in each description
- End each description with a soft call to action
- Return a table: option, title, title characters, description, description characters

What to check: The model counts characters unreliably - paste the winner into a pixel-width checker before shipping, and confirm every claim is actually true on the page.

In this lab: Run it on the EcoDesk product page. Every claim must trace to visible page content; EcoDesk has no published price, rating, certification or environmental result to cite.

### Template 2 - Bulk metas for a whole section of the site

Where AI genuinely changes the economics: a hundred course or product pages that would take a week by hand take an afternoon with review.

You are an SEO assistant. I will paste a table of URLs with columns: URL, page type, primary keyword, current meta title.

For each row write a new meta title (max 60 chars) and meta description (max 155 chars).

Rules:

- Keep each row's primary keyword in the first half of the title
- Never reuse the same opener twice - vary the phrasing across rows
- Append " | [BRAND]" only where it still fits the limit
- Flag any two rows that target the same query (cannibalisation risk)
- Return one table I can paste into a spreadsheet

[PASTE YOUR TABLE]

What to check: Spot-check ten rows, not two. Bulk runs fail in patterns - if the model drifts into a repeated formula halfway down, regenerate from that row with "vary the openers" restated.

In this lab: Paste the EcoDesk page inventory (index, products, guide, article, thank-you). The cannibalisation flag is the graded output: two EcoDesk pages chasing one query is a clustering defect, not a copy defect.

## Pre-publish checklist

Whatever the prompt produced, it goes through the same 60-second review before it ships.

| Check                                             | Why it matters                                                                          |
|---------------------------------------------------|-----------------------------------------------------------------------------------------|
| Character/pixel limits verified in a real checker | Models miscount; Google truncates at pixel width, not characters.                       |
| Every number and claim confirmed true             | Fabricated specifics are the most common AI failure - and the most damaging.            |
| Primary keyword present and early                 | The one on-page signal a meta rewrite must never lose.                                  |
| Reads differently from your last five outputs     | Repeated AI phrasing across pages and posts is what "thin content" looks like at scale. |
| A named human edited it                           | Generation is a drafting step, not a publishing pipeline.                               |

These templates are model-agnostic: the same wording works in ChatGPT, Claude, Gemini and Copilot.

## Prompt usage record

| Field          | Record                    |
|----------------|---------------------------|
| Tool/model     |                           |
| Date/time      |                           |
| Prompt version | v15.0                     |
| Input sources  |                           |
| Reviewer       |                           |
| Decision       | Approve / Revise / Reject |
| Human changes  |                           |