

Prompt Pack - Lab 5: Create a People-First SEO Blog Article

Replace bracketed fields with verified inputs. Never paste personal, confidential, licensed or credential data into a tool without authority.

Prompt 1 - CONTEXT task prompt

Context: You support EcoDesk, a fictional Singapore sustainable-office-products business. Audience: [role]. Page/task: [page]. Brand facts: [verified facts only].

Objective: [one measurable outcome].

Necessary inputs: [paste source text/data with source and date].

Task steps: Draft an answer-first article one section at a time using only the approved claim ledger. Mark missing evidence as [VERIFY] and create no environmental claims.

Evidence: Cite the exact supplied row, page element or source for every material conclusion. Separate observation, inference and recommendation.

eXclusions: No invented metrics, rankings, citations, testimonials, prices, certifications, environmental benefits or legal claims. No keyword stuffing, link schemes or private data.

Template: Return findings/options, evidence, uncertainty, recommendation, human checks and acceptance checklist.

Prompt 2 - Adversarial quality review

Act as a critical SEO editor. Review intent, accuracy, originality, clarity, brand, accessibility, Google Search policies and the lab acceptance test. Return issue, severity, evidence, correction and reviewer decision.

Prompt 3 - Human approval summary

Prepare a change preview. Show affected files/fields, before/after meaning, evidence, limitations, tests, risks, rollback and an approve/revise/reject recommendation. Do not publish.

Field-tested prompt templates

Published templates from [AI Prompts for SEO Meta Tags & LinkedIn Posts \(Templates\)](#). Distribution copy is drafted from the finished article, using only claims the article already supports.

Every template follows the same five-part skeleton - it is the skeleton, not the tool, that does the work:

Part	What it does
Role	Tell the model what it is ("you are an SEO copywriter"), which sets vocabulary and priorities.
Context	The page, audience, keyword and differentiator. The model cannot know your business; feed it.
Constraints	Character limits, banned phrases, tone rules. Constraints are what separate usable output from slop.
Format	Ask for a table or labelled options so the output drops straight into your workflow.
Guardrail	An explicit instruction not to invent facts, and to ask when information is missing.

Replace the [BRACKETED] fields with verified inputs before running any template below.

Template 1 - Turn an article into a LinkedIn post

The most common repurposing job. The two rules that matter most: the hook must survive the "...see more" fold, and the post must contain one concrete number so it reads like experience rather than content marketing.

Turn the article below into a LinkedIn post for my company page.

Structure:

- Hook: the first 2 lines must work before the "...see more" fold - no throat-clearing
- Body: 3-5 short paragraphs, one idea each, plain English
- Include one concrete example or number from the article
- CTA: one line inviting readers to the full article at [URL]
- 3-5 hashtags: one broad, the rest niche

Rules: no emojis in the first line, never open with "I'm excited to share", 150-220 words total.

[PASTE ARTICLE OR KEY POINTS]

What to check: The concrete number must come from the article, not from the model. If the article has no number, the honest fix is to add one to the article first.

In this lab: Feed it the EcoDesk guide page. The post must seek earned attention only - no reciprocal-link request, no paid placement, no engagement pod.

Pre-publish checklist

Whatever the prompt produced, it goes through the same 60-second review before it ships.

Check	Why it matters
Character/pixel limits verified in a real checker	Models miscount; Google truncates at pixel width, not characters.
Every number and claim confirmed true	Fabricated specifics are the most common AI failure - and the most damaging.
Primary keyword present and early	The one on-page signal a meta rewrite must never lose.
Reads differently from your last five outputs	Repeated AI phrasing across pages and posts is what "thin content" looks like at scale.
A named human edited it	Generation is a drafting step, not a publishing pipeline.

These templates are model-agnostic: the same wording works in ChatGPT, Claude, Gemini and Copilot.

Prompt usage record

Field	Record
Tool/model	
Date/time	
Prompt version	v15.0
Input sources	

Field	Record
Reviewer	
Decision	Approve / Revise / Reject
Human changes	